

Assignment / Explore Query Planning

Connor Clancy - clancy.co@northeastern.edu

Spring 2023

Library Imports

```
library(RSQLite)
library(RMySQL)

## Loading required package: DBI
##
## Attaching package: 'RMySQL'
## The following object is masked from 'package:RSQLite':
##
##      isIdCurrent

library(readr)
library(sqldf)

## Loading required package: gsubfn
## Loading required package: proto
## sqldf will default to using MySQL
options(sqldf.driver = "SQLite")
```

Connect to SQLite Database

```
fpath = getwd()
dbfile = "/sakila.db"

# connect to the database if exists, else create a new database
lcon <- dbConnect(RSQLite::SQLite(), paste0(fpath, dbfile))

dbGetQuery(lcon, "SELECT * FROM film LIMIT 5;")
```

```
##      film_id      title
## 1          1 ACADEMY DINOSAUR
## 2          2  ACE GOLDFINGER
## 3          3 ADAPTATION HOLES
## 4          4 AFFAIR PREJUDICE
## 5          5   AFRICAN EGG
##
## 1                      A Epic Drama of a Feminist And a Mad Scientist who must Battle a Teacher in T
## 2          A Astounding Epistle of a Database Administrator And a Explorer who must Find a C
```

```

## 3          A Astounding Reflection of a Lumberjack And a Car who must Sink a Lumberjack :
## 4          A Fanciful Documentary of a Frisbee And a Lumberjack who must Chase a Mon
## 5 A Fast-Paced Documentary of a Pastry Chef And a Dentist who must Pursue a Forensic Psychologist in
##   release_year language_id original_language_id rental_duration rental_rate
## 1      2006          1          NA              6          0.99
## 2      2006          1          NA              3          4.99
## 3      2006          1          NA              7          2.99
## 4      2006          1          NA              5          2.99
## 5      2006          1          NA              6          2.99
##   length replacement_cost rating          special_features
## 1     86          20.99      PG Deleted Scenes,Behind the Scenes
## 2     48          12.99       G      Trailers,Deleted Scenes
## 3     50          18.99  NC-17      Trailers,Deleted Scenes
## 4    117          26.99       G Commentaries,Behind the Scenes
## 5    130          22.99       G      Deleted Scenes
##           last_update
## 1 2006-02-15 05:03:42
## 2 2006-02-15 05:03:42
## 3 2006-02-15 05:03:42
## 4 2006-02-15 05:03:42
## 5 2006-02-15 05:03:42

```

Connect to MySQL Database

```

db_user <- 'admin'
db_password <- 'Northea$tern23'
db_name <- 'sakila'
db_host <- 'cclancy-cs5200.cbowkysg1oyc.us-east-2.rds.amazonaws.com'
db_port <- 3306
mscon <- dbConnect(MySQL(), user = db_user, password = db_password,
                    dbname = db_name, host = db_host, port = db_port)

```

```

dbGetQuery(mscon, "SELECT * FROM film LIMIT 5;")

```

```

## Warning in .local(conn, statement, ...): Unsigned INTEGER in col 0 imported as
## numeric
## Warning in .local(conn, statement, ...): Unsigned INTEGER in col 4 imported as
## numeric
## Warning in .local(conn, statement, ...): Unsigned INTEGER in col 5 imported as
## numeric
## Warning in .local(conn, statement, ...): Unsigned INTEGER in col 6 imported as
## numeric
## Warning in .local(conn, statement, ...): Decimal MySQL column 7 imported as
## numeric
## Warning in .local(conn, statement, ...): Unsigned INTEGER in col 8 imported as
## numeric
## Warning in .local(conn, statement, ...): Decimal MySQL column 9 imported as
## numeric
## Warning in .local(conn, statement, ...): unrecognized MySQL field type 7 in
## column 12 imported as character

```

```

##      film_id      title
## 1         1 ACADEMY DINOSAUR
## 2         2   ACE GOLDFINGER
## 3         3 ADAPTATION HOLES
## 4         4 AFFAIR PREJUDICE
## 5         5   AFRICAN EGG
##
## 1          A Epic Drama of a Feminist And a Mad Scientist who must Battle a Teacher in T
## 2          A Astounding Epistle of a Database Administrator And a Explorer who must Find a C
## 3          A Astounding Reflection of a Lumberjack And a Car who must Sink a Lumberjack
## 4          A Fanciful Documentary of a Frisbee And a Lumberjack who must Chase a Mon
## 5 A Fast-Paced Documentary of a Pastry Chef And a Dentist who must Pursue a Forensic Psychologist in
##      release_year language_id original_language_id rental_duration rental_rate
## 1         2006         1             NA             6             0.99
## 2         2006         1             NA             3             4.99
## 3         2006         1             NA             7             2.99
## 4         2006         1             NA             5             2.99
## 5         2006         1             NA             6             2.99
##      length replacement_cost rating      special_features
## 1        86          20.99     PG Deleted Scenes,Behind the Scenes
## 2        48          12.99      G      Trailers,Deleted Scenes
## 3        50          18.99  NC-17      Trailers,Deleted Scenes
## 4       117          26.99      G Commentaries,Behind the Scenes
## 5       130          22.99      G      Deleted Scenes
##
##      last_update
## 1 2006-02-15 05:03:42
## 2 2006-02-15 05:03:42
## 3 2006-02-15 05:03:42
## 4 2006-02-15 05:03:42
## 5 2006-02-15 05:03:42

```

Tasks

Question 1

Ensuring that no user-defined indexes exist (delete all user-defined indexes, if there are any), find the number of films per category. The query should return the category name and the number of films in each category. Show us the code that determines if there are any indexes and the code to delete them if there are any.

```
delete_idx <- function (tables, dbcon, database) {  
  
  if (database == "mysql") {  
    for (t in tables) {  
      df_idx <- dbGetQuery(dbcon, sprintf("SHOW INDEXES FROM %s  
                                          WHERE Key_name != 'PRIMARY'  
                                          AND Key_name NOT LIKE '%%fk%%';", t))  
  
      for (i in df_idx$Key_name) {  
        dbExecute(dbcon, sprintf("DROP INDEX %s ON %s;", i, t))  
      }  
    }  
  } else if (database == "sqlite") {  
    for (t in tables) {  
      df_idx <- dbGetQuery(dbcon, sprintf("SELECT name FROM sqlite_master  
                                          WHERE type == 'index'  
                                          AND tbl_name == '%s'  
                                          AND name NOT LIKE '%%autoindex%%'", t))  
  
      for (i in df_idx$name) {  
        dbExecute(dbcon, sprintf("DROP INDEX %s", i))  
      }  
    }  
  }  
}
```

```
delete_idx(c('film', 'category', 'film_category'), lcon, "sqlite")
```

```
dbGetQuery(lcon, "  
  SELECT  
    c.category_id,  
    c.name AS category_name,  
    COUNT(f.film_id) AS film_count  
  FROM film AS f  
  INNER JOIN film_category AS j  
    ON f.film_id = j.film_id  
  INNER JOIN category AS c  
    ON c.category_id = j.category_id  
  GROUP BY  
    c.category_id,  
    c.name  
  ORDER BY  
    COUNT(f.film_id) DESC;  
")
```

```
##      category_id category_name film_count
```

## 1	15	Sports	74
## 2	9	Foreign	73
## 3	8	Family	69
## 4	6	Documentary	68
## 5	2	Animation	66
## 6	1	Action	64
## 7	13	New	63
## 8	7	Drama	62
## 9	10	Games	61
## 10	14	Sci-Fi	61
## 11	3	Children	60
## 12	5	Comedy	58
## 13	4	Classics	57
## 14	16	Travel	57
## 15	11	Horror	56
## 16	12	Music	51

Question 2

Ensuring that no user-defined indexes exist (delete all user-defined indexes, if there are any), execute the same query (same SQL) as in (1) but against the MySQL database. Make sure you reuse the same SQL query string as in (1)

```
delete_idx(c('film', 'category', 'film_category'), mscon, "mysql")
```

```
## Warning in .local(conn, statement, ...): Unsigned INTEGER in col 3 imported as
## numeric
```

```
## Warning in .local(conn, statement, ...): unrecognized MySQL field type 6 in
## column 8 imported as character
```

```
## Warning in .local(conn, statement, ...): Unsigned INTEGER in col 3 imported as
## numeric
```

```
## Warning in .local(conn, statement, ...): unrecognized MySQL field type 6 in
## column 8 imported as character
```

```
## Warning in .local(conn, statement, ...): Unsigned INTEGER in col 3 imported as
## numeric
```

```
## Warning in .local(conn, statement, ...): unrecognized MySQL field type 6 in
## column 8 imported as character
```

```
dbGetQuery(mscon, "
SELECT
  c.category_id,
  c.name AS category_name,
  COUNT(f.film_id) AS film_count
FROM film AS f
INNER JOIN film_category AS j
  ON f.film_id = j.film_id
INNER JOIN category AS c
  ON c.category_id = j.category_id
GROUP BY
  c.category_id,
  c.name
ORDER BY
```

```

COUNT(f.film_id) DESC;
")

```

```

## Warning in .local(conn, statement, ...): Unsigned INTEGER in col 0 imported as
## numeric

```

```

##   category_id category_name film_count
## 1          15         Sports          74
## 2           9        Foreign          73
## 3           8         Family          69
## 4           6    Documentary          68
## 5           2        Animation          66
## 6           1         Action          64
## 7          13           New          63
## 8           7         Drama          62
## 9          10         Games          61
## 10          14        Sci-Fi          61
## 11           3       Children          60
## 12           5         Comedy          58
## 13           4        Classics          57
## 14          16         Travel          57
## 15          11         Horror          56
## 16          12          Music          51

```

Question 3

Find out how to get the query plans for SQLite and MySQL and then display the query plans for each of the query executions in (1) and (2).

```

dbGetQuery(lconn, "
    EXPLAIN QUERY PLAN SELECT
        c.category_id,
        c.name AS category_name,
        COUNT(f.film_id) AS film_count
    FROM film AS f
    INNER JOIN film_category AS j
        ON f.film_id = j.film_id
    INNER JOIN category AS c
        ON c.category_id = j.category_id
    GROUP BY
        c.category_id,
        c.name
    ORDER BY
        COUNT(f.film_id) DESC;
")

```

```

##   id parent notused
## 1  9      0        0
## 2 11      0        0
## 3 14      0        0
## 4 17      0        0
## 5 57      0        0
##
##                                     detail
## 1 SCAN j USING COVERING INDEX sqlite_autoindex_film_category_1
## 2          SEARCH f USING INTEGER PRIMARY KEY (rowid=?)

```

```
## 3          SEARCH c USING INTEGER PRIMARY KEY (rowid=?)
## 4                                USE TEMP B-TREE FOR GROUP BY
## 5                                USE TEMP B-TREE FOR ORDER BY
```

```
dbGetQuery(mscon, "
    EXPLAIN SELECT
      c.category_id,
      c.name AS category_name,
      COUNT(f.film_id) AS film_count
    FROM film AS f
    INNER JOIN film_category AS j
      ON f.film_id = j.film_id
    INNER JOIN category AS c
      ON c.category_id = j.category_id
    GROUP BY
      c.category_id,
      c.name
    ORDER BY
      COUNT(f.film_id) DESC;
  ")
```

```
##  id select_type table partitions  type                possible_keys
## 1  1      SIMPLE   c      <NA>      ALL                PRIMARY
## 2  1      SIMPLE   j      <NA>      ref PRIMARY,fk_film_category_category
## 3  1      SIMPLE   f      <NA>      eq_ref            PRIMARY
##                                key key_len          ref rows filtered
## 1                                <NA>      <NA>          <NA>      16      100
## 2 fk_film_category_category      1 sakila.c.category_id 62      100
## 3                                PRIMARY      2      sakila.j.film_id  1      100
##                                Extra
## 1 Using temporary; Using filesort
## 2                                Using index
## 3                                Using index
```

Question 4

Comment on the differences between the query plans? Are they the same? How do they differ? Why do you think they differ? Do both take the same amount of time?

The two plans are not the same. SQLite appears to have extra steps for the `GROUP BY` and `ORDER BY` clauses that MySQL is able to handle within one of the other steps in the plan. It appears to me that MySQL runs quicker.

Question 5

Write a SQL query against the SQLite database that returns the title, language and length of the film with the title “ZORRO ARK”.

```
dbGetQuery(lcon, "
    SELECT
      f.title,
      l.name AS language,
      f.length
    FROM film AS f
    INNER JOIN language AS l
      ON f.language_id = l.language_id
```

```
WHERE title = 'ZORRO ARK';
")
```

```
##      title language length
## 1 ZORRO ARK   English    50
```

Question 6

For the query in (5), display the query plan.

```
dbGetQuery(lcon, "
    EXPLAIN QUERY PLAN SELECT
        f.title,
        l.name AS language,
        f.length
    FROM film AS f
    INNER JOIN language AS l
        ON f.language_id = l.language_id
    WHERE title = 'ZORRO ARK';
")
```

```
##   id parent notused      detail
## 1  3      0        0      SCAN f
## 2  7      0        0 SEARCH 1 USING INTEGER PRIMARY KEY (rowid=?)
```

Question 7

In the SQLite database, create a user-defined index called “TitleIndex” on the column TITLE in the table FILM.

```
dbExecute(lcon,
    'CREATE UNIQUE INDEX IF NOT EXISTS TitleIndex
    ON film (title);')
```

```
## [1] 0
```

Question 8

Re-run the query from (5) now that you have an index and display the query plan.

```
dbGetQuery(lcon, "
    EXPLAIN QUERY PLAN SELECT
        f.title,
        l.name AS language,
        f.length
    FROM film AS f
    INNER JOIN language AS l
        ON f.language_id = l.language_id
    WHERE title = 'ZORRO ARK';
")
```

```
##   id parent notused      detail
## 1  4      0        0 SEARCH f USING INDEX TitleIndex (title=?)
## 2  9      0        0 SEARCH 1 USING INTEGER PRIMARY KEY (rowid=?)
```


Question 9

Are the query plans the same in (6) and (8)? What are the differences? Is there a difference in execution time? How do you know from the query plan whether it uses an index or not?

No, the query plans are not the same. In (6) the query scans the whole table to find the movie, where as in (8) the query plan states that it is going to use the index created to complete the search. The query with the index performs faster.

Question 10

Write a SQL query against the SQLite database that returns the title, language and length of all films with the word “GOLD” with any capitalization in its name, i.e., it should return “Gold Finger”, “GOLD FINGER”, “THE GOLD FINGER”, “Pure GOLD” (these are not actual titles).

```
dbGetQuery(lcon, "
    SELECT
        f.title,
        l.name AS language,
        f.length
    FROM film AS f
    INNER JOIN language AS l
        ON f.language_id = l.language_id
    WHERE lower(title) LIKE '%gold%';
")
```

##		title	language	length
## 1	ACE	GOLDFINGER	English	48
## 2	BREAKFAST	GOLDFINGER	English	123
## 3		GOLD RIVER	English	154
## 4	GOLDFINGER	SENSIBILITY	English	93
## 5		GOLDMINE TYCOON	English	153
## 6		OSCAR GOLD	English	115
## 7	SILVERADO	GOLDFINGER	English	74
## 8		SWARM GOLD	English	123

Question 11

Get the query plan for (10). Does it use the index you created? If not, why do you think it didn't?

```
dbGetQuery(lcon, "
    EXPLAIN QUERY PLAN SELECT
        f.title,
        l.name AS language,
        f.length
    FROM film AS f
    INNER JOIN language AS l
        ON f.language_id = l.language_id
    WHERE lower(title) LIKE '%gold%';
")
```

##	id	parent	notused	detail
## 1	3	0	0	SCAN f
## 2	9	0	0	SEARCH 1 USING INTEGER PRIMARY KEY (rowid=?)

The query does not use the Index we created because we are using a wild card to search through the column. Since we have to search every single row for the wildcard string to see if it contains 'gold' SQLite appears to just choose to scan the whole table directly rather than using the index.